

Audit of a private-market growth-company prospectus

*Anonymised from a 2019 Form S-1 that was withdrawn 33 days after filing.
Independently re-verifiable hashed evidence record produced by the Decision Intel pipeline.*

— §1 INTEGRITY FINGERPRINTS

What this record is and how to confirm it

The cryptographic provenance below is the audit committee's first read. Every field is independently verifiable; the verification URL below resolves to the same hashes.

Record id	dpr_202604230900_b341249d
Audit timestamp	2026-04-23 09:00:00 UTC
Pipeline version	di-pipeline · 10 nodes · v2.1.0
Input document hash	b341249dae4ee8e4...3c748f90 SHA-256
Prompt fingerprint	5c9f71f450af627d...4b5c3abd SHA-256
Record schema	v2.1.0 FORWARD-COMPATIBLE
Tamper-evidence	SHA-256 input hash + record fingerprint
Private-key signing	Planned · Q3 2026 roadmap

🔗 HOW TO VERIFY THIS RECORD

Re-hash the source memo with SHA - 256 and compare to the input - document hash above. The prompt fingerprint is the SHA - 256 of the prompt version active at audit time — a divergent hash on a re-run proves the prompts evolved between audits, which is itself recorded against this record id. Any byte change in the source memo invalidates the input hash above; the failure surfaces at re-hash time.

```
decision-intel.com/verify/dpr_202604230900_b341249d
```

Tamper-evidence today: SHA - 256 over canonicalised inputs. Private-key signing of the record itself is on the published roadmap (Q3 2026); the schema is forward-compatible. The record fingerprint will become a digital signature without breaking re-verification of older records.

DISCLOSURE Decision Intel runs on SOC 2 Type II infrastructure (Vercel + Supabase) and is aligned with the eleven internationally-recognised AI governance principles codified by AI Verify (Singapore IMDA). Decision Intel's own product-level SOC 2 Type I audit is targeted for Q4 2026; AI Verify alignment is self-attested and no third-party audit has yet been performed against the mapping. Vocabulary in this record is calibrated to be defensible during a vendor-risk register review.

§2 METHODOLOGICAL DEFENSIBILITY

Judge variance — what the jury saw

Decision Intel runs three independent professional lenses on every memo: equity -research skeptical, regulator -hostile, and contrarian -strategist. Below: the variance the jury produced. Low variance means the audit verdict is robust across lenses; high variance tells the reviewer which audience will be harshest.

MEAN NOISE SCORE

22/100

Composite across three R²F lenses.

STD. DEVIATION

—

Variance across the jury sample. Lower = stronger convergence.

INTER - JUDGE VARIANCE

—

Variance not recorded for legacy audits.

BIAS DETECTIVE FLAGS

—

Single independent pass; cross-checked against the noise jury.

META - JUDGE VERDICT *Concurred on three of four material flags; minority dissent on reference -class (Pre -Mortem flagged the 18 -month break -even claim as optimistic; Ensemble Sampling concurred only after base -rate pull).*

ON PER - JUDGE IP Summary view of judge variance. Per -judge granular outputs are stored in the internal audit log and are available on request under the DPA; they are deliberately excluded from the client -facing record to protect prompt internals.

§3 REPRODUCIBILITY

Model lineage and prompt fingerprint

The exact configuration in force at audit time. Together with the prompt -fingerprint hash on page 1, this establishes that any re -run of this audit on the same source memo would produce a fingerprint identical to the one on file — or fail audibly.

Gdpr Anonymizer	preprocessing-tier T=0.00 · top-p=0.95
Data Structurer	preprocessing-tier T=0.00 · top-p=0.95
Intelligence Gatherer	preprocessing-tier T=0.20 · top-p=0.95
Bias Detective	analysis-tier T=0.20 · top-p=0.95
Noise Judge	analysis-tier T=0.40 · top-p=0.95
Statistical Jury	analysis-tier T=0.30 · top-p=0.95
Rpd Recognition	analysis-tier T=0.25 · top-p=0.95
Forgotten Questions	analysis-tier T=0.35 · top-p=0.95

DECISION PROVENANCE RECORDPAGE 3 OF 7DECISION-INTEL.COM

Meta Judge

analysis-tier T=0.15 · top-p=0.95

Risk Scorer

deterministic T=0.00 · top-p=1.00

ON COST - TIER ROUTING Cost-tier routing: preprocessing and analysis tiers run on separate model classes; final risk score is deterministic, not model-generated. Actual model IDs are resolved at audit time and available to design partners on request under NDA.

§4 RECOGNITION-RIGOR FRAMEWORK

Procurement-grade audit signals

The five strips below operationalise Kahneman & Klein (2009) "Conditions for Intuitive Expertise." Each strip independently answers a question a procurement reviewer should be able to ask of any AI-augmented decision audit. When a signal cannot be honestly produced (insufficient data, novel decision class), the strip is omitted rather than fabricated.

§4.1 · VALIDITY

LOW VALIDITY ENVIRONMENT

Low-validity domain — narrative coherence does NOT confer credibility.
Methodology v2.1.0 reweighs toward historical alignment + bias load and away from evidence quality.

documentType="ic_memo" maps to low-validity per the Kahneman & Klein 2009 environment taxonomy; industry="real_estate" tilts the band down (already at low). Per Kahneman & Klein (2009), validity is the first precondition for trustworthy intuition; in low- and zero-validity environments the audit methodology shifts weights to compensate for the documented unreliability of pattern recognition in those domains.

Document type	ic_memo
Industry	real_estate
Decision horizon	18-month integration horizon

§4.2 · REFERENCE CLASS FORECAST

CHALLENGING BASE RATE

11 historically-similar decisions; 64% failed.

Reference class of 11 historically-similar decisions failed in 64% of cases. Per Kahneman & Lovallo (2003), this is a structurally challenging base rate — the memo's confidence should be calibrated against this rate, not against the inside-view narrative. Closest analog: WeWork (2019) — outcome: catastrophic failure. Per Kahneman & Lovallo (2003) "Delusions of Success," the inside-view narrative ALWAYS feels more compelling than the outside-view base rate — but the base rate wins on average. The memo's confidence should be calibrated against the matched-class failure rate, not against the inside-view story.

Pool searched	143 cases · 11 matched
Top analog	The WeWork S-1 Collapse (2019) · sim 74%
Baseline sample	11 closed cases

Few closed-loop outcomes — experience-based claims should be cross-checked against external base rates.

Only 2 closed-loop outcomes on market_entry decisions in the past 18 months — too few for calibrated intuition by Kahneman & Klein's (2009) standard. Experience-based justifications in this memo should be cross-checked against external base rates. Mean Brier 0.210 across 4 scored outcomes (CIA-analyst band). Per Kahneman & Klein (2009), the second precondition for trustworthy intuition is adequate opportunity to learn from rapid feedback. Domains with sparse feedback (M&A, market entry, long-horizon strategy) require independent base-rate cross-checks even when the operator is highly experienced.

Closed outcomes (lifetime)	6
Recent (18 months)	4
Mean Brier	0.210
Days since last outcome	95

§4.4 · ORG CALIBRATION

28 OUTCOMES · GOOD BAND

Calibrated against 28 closed decisions for this organisation · mean Brier 0.183 · recalibrated -4 from absolute.

DQI shown is calibrated against this organisation's outcome history (28 closed decisions, mean Brier 0.183 · good). Original-vs-recalibrated delta: -4 (the organisation has historically been over-confident on DACH-style market entries; the recalibrator pulls the headline score down accordingly). The DQI shown on this audit reflects this organisation's calibrated-by-outcome history; per-org calibration is the platform's structural answer to the Cloverpop data-advantage attack vector.

Decisions tracked	42
Outcomes closed	28
Mean Brier	0.183
Calibration band	good

§4.5 · COUNTERFACTUAL IMPACT

+13.1% IF ALL FLAGGED BIASES ADDRESSED

Top scenario: addressing Overconfidence Bias alone would have lifted the verdict by 18.4%.

Monetary anchors derived from the linked DecisionFrame.monetaryValue (DACH expansion budget USD 14M). Confidence reflects historical sample size (Wilson score) weighted by per-org CausalEdge strength. Aggregate assumes bias independence — the 32.4-pt headline overstates the realistic combined effect; reviewers should trust the confidence-weighted 13.1-pt number for board-presentation purposes.

Overconfidence Bias	+18.4% · n=31
Confirmation Bias	+11.2% · n=26
Sunk Cost Fallacy	+7.6% · n=19

ON VOCABULARY The R²F is the platform's integration of Kahneman's rigour (System 2 debiasing) and Klein's recognition (System 1 amplification), arbitrated in one pipeline. Anchor citation: Kahneman & Klein (2009) "Conditions for Intuitive Expertise: a failure to disagree." Methodology version active at audit time is recorded with the prompt fingerprint on page 1 — divergent re-runs surface as fingerprint mismatches.